

Artificial Intelligence



Hai Thi Tuyet Nguyen

Outline

CHAPTER 1: INTRODUCTION (CHAPTER 1)

CHAPTER 2: INTELLIGENT AGENTS (CHAPTER 2)

CHAPTER 3: SOLVING PROBLEMS BY SEARCHING (CHAPTER 3)

CHAPTER 4: INFORMED SEARCH (CHAPTER 3)

CHAPTER 5: LOGICAL AGENT (CHAPTER 7)

CHAPTER 6: FIRST-ORDER LOGIC (CHAPTER 8, 9)

CHAPTER 7: QUANTIFYING UNCERTAINTY (CHAPTER 13)

CHAPTER 8: PROBABILISTIC REASONING (CHAPTER 14)

CHAPTER 9: LEARNING FROM EXAMPLES (CHAPTER 18)

CHAPTER 5:

LOGICAL AGENT

- 5.1 Knowledge-Based Agents
- 5.2 The Wumpus World
- 5.3 Logic
- 5.4 Propositional Logic
- 5.5 Propositional Theorem Proving
- 5.6 Inference Rules, Theorem Proving

5.1 Knowledge-Based Agents

- *Knowledge base (KB)* = a set of sentences in a *formal* language (i.e., knowledge representation language)
- *Declarative* approach to build an agent:
 - *TELL* it what it needs to know
 - *ASK* itself what to do - answer should follow from the KB

5.1 A simple knowledge-based agent.

- The agent takes a percept as input and returns an action.
It maintains a knowledge base
- How it works:
 - TELLS the knowledge base what it perceives.
 - ASKS the knowledge base what action it should perform.
 - TELLS the knowledge base which action was chosen, and executes the action.

function **KB-AGENT**(*percept*) **returns** an *action*

static: *KB*, a knowledge base

t, a counter, initially 0, indicating time

TELL(*KB*, MAKE-PERCEPT-SENTENCE(*percept*, *t*))

action ← ASK(*KB*, MAKE-ACTION-QUERY(*t*))

TELL(*KB*, MAKE-ACTION-SENTENCE(*action*, *t*))

t ← *t* + 1

return *action*

5.2 Wumpus World PEAS description

- **Performance measure:**
 - +1000: climb out of the cave with the gold,
 - -1000: fall into a pit or being eaten by the wumpus
 - -1: each action
 - -10: use up the arrow.
 - The game ends: the agent dies or it climbs out of the cave
- **Environment:** A 4×4 grid of rooms.
 - Start location of the agent: the square labeled [1,1]
 - Locations of the gold and the wumpus: random
- **Actuators:** Move Forward, Turn Left, Turn Right, Grab, Climb, Shoot
- **Sensors:** the agent will perceive
 - Stench: in the square containing the monster (called wumpus) and in the directly adjacent squares
 - Breeze: in the squares directly adjacent to a pit
 - Glitter: in the square where the gold is
 - Bump: into a wall
 - Scream: anywhere in the cave when the wumpus is killed


5.2 Wumpus World PEAS description

- The first percept is [None, None, None, None, None] => the agent can conclude that its neighboring squares, [1,2] and [2,1], are OK.

OK			
OK <div>A</div>	OK		


5.2 Wumpus World PEAS description

The agent decides to move forward to [2,1].




5.2 Wumpus World PEAS description

- The agent perceives a breeze (denoted by “B”) in $[2,1]$ \Rightarrow there must be a pit in a neighboring square.
- The pit cannot be in $[1,1]$ \Rightarrow so there must be a pit in $[2,2]$ or $[3,1]$ or both.



5.2 Wumpus World PEAS description


The agent will turn around, go back to [1,1], and then proceed to [1,2].



5.2 Wumpus World PEAS description


- The agent perceives a stench in $[1,2] \Rightarrow$ there must be a wumpus nearby ($[2,2]$ or $[1,3]$)
- The lack of stench when the agent was in $[2,1] \Rightarrow$ wumpus cannot be in $[2,2] \Rightarrow$ it is in $[1,3]$
- the lack of a breeze in $[1,2] \Rightarrow$ there is no pit in $[2,2]$

$\Rightarrow [2,2]$: safe, OK



5.2 Wumpus World PEAS description

The agent draws a conclusion from the available information, that conclusion is guaranteed to be correct if the available information is correct.



5.3 Logic

- Logics are formal languages for representing information

- *Syntax* defines the sentences in the language

E.g., “ $x + y = 4$ ” is a well-formed sentence, whereas “ $x4y+ =$ ” is not

- *Semantics* defines the “meaning” of sentences

- The semantics defines the *truth* of each sentence with respect to each *possible world* (i.e., *model*).

E.g., the sentence “ $x + y = 4$ ” is true in a world where x is 2 and y is 2,
but false in a world where x is 1 and y is 1


5.3 Entailment

- Entailment means that *one thing follows from another*:
- Knowledge base KB entails sentence α if and only if α is true in all worlds where KB is true
 $KB \models \alpha$
- E.g., KB containing “the Giants won” and “the Reds won” entails “Either the Giants won or the Reds won”

5.3 Models

- **Models** are formally structured worlds with respect to which truth can be evaluated
- We say m is a **model** of a sentence α if α is true in m
- $M(\alpha)$ is the set of all models of α
- Then $KB \models \alpha$ if and only if $M(KB) \subseteq M(\alpha)$

E.g. $KB = \text{Giants won and Reds won}$
 $\alpha = \text{Giants won}$



5.3 Inference and Entailment

- An inference algorithm is a procedure for deriving a sentence from the KB
- If an inference algorithm i can derive α from KB, we write

$$KB \vdash_i \alpha$$

which is pronounced “ α is derived from KB by i ” or “ i derives α from KB”

OR: the sentence α is inferred from KB using algorithm i .

5.4 Propositional logic: Syntax

- Propositional logic is the simplest logic - illustrates basic ideas
- The proposition symbols P_1, P_2, \dots are sentences
- If S is a sentence, $\neg S$ is a sentence (negation)
- If S_1 and S_2 are sentences, $S_1 \wedge S_2$ is a sentence (conjunction)
- If S_1 and S_2 are sentences, $S_1 \vee S_2$ is a sentence (disjunction)
- If S_1 and S_2 are sentences, $S_1 \Rightarrow S_2$ is a sentence (implication)
- If S_1 and S_2 are sentences, $S_1 \Leftrightarrow S_2$ is a sentence (biconditional)

5.4 Propositional logic: Semantics

- Each model specifies true/false for each proposition symbol

E.g. $P_{1,2}$ $P_{2,2}$ $P_{3,1}$
true true false

- Rules for evaluating truth with respect to a model m :

$\neg S$	is true iff	S	is false	
$S_1 \wedge S_2$	is true iff	S_1	is true and	S_2 is true
$S_1 \vee S_2$	is true iff	S_1	is true or	S_2 is true
$S_1 \Rightarrow S_2$	is true iff	S_1	is false or	S_2 is true
	i.e., is false iff	S_1	is true and	S_2 is false
$S_1 \Leftrightarrow S_2$	is true iff	$S_1 \Rightarrow S_2$	is true and	$S_2 \Rightarrow S_1$ is true

5.4 Truth tables for connectives

P	Q	$\neg P$	$P \wedge Q$	$P \vee Q$	$P \Rightarrow Q$	$P \Leftrightarrow Q$
false	false	true	false	false	true	true
false	true	true	false	true	true	false
true	false	false	false	true	false	false
true	true	false	true	true	true	true

5.4 A simple knowledge base - Wumpus world sentences

$P_{x,y}$ is true if there is a pit in $[x, y]$.

$B_{x,y}$ is true if the agent perceives a breeze in $[x, y]$.

There is no pit in $[1,1]$: $R_1 : \neg P_{1,1}$.

A square is breezy if and only if there is a pit in a neighboring square.


$$\text{R2: } B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1}).$$

$$\text{R3 : } B_{2,1} \Leftrightarrow (P_{1,1} \vee P_{2,2} \vee P_{3,1}).$$

The breeze percepts on the first two squares of the agent

$$\text{R4} : \neg B_{1,1}.$$

R5 : $B_{2,1}$.



5.4 A simple knowledge base - Wumpus world sentences

- Goal: to decide whether $KB \models \alpha$ for some sentence α
 KB, α as $\neg P_{1,2}$
prove: $KB \models \neg P_{1,2}$
- A model-checking approach:
 - enumerate the models
 - check that α is true in every model in which KB is true

5.4 A simple knowledge base - Wumpus world sentences

With 7 symbols, there are $2^7 = 128$ possible models; in 3 of these, KB is true.

In those 3 models, $\neg P_{1,2}$ is true or there is no pit in [1,2].

$B_{1,1}$	$B_{2,1}$	$P_{1,1}$	$P_{1,2}$	$P_{2,1}$	$P_{2,2}$	$P_{3,1}$	R_1	R_2	R_3	R_4	R_5	KB
false	false	false	false	false	false	false	true	true	true	true	false	false
false	false	false	false	false	false	true	true	true	false	true	false	false
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮
false	true	false	false	false	false	false	true	true	false	true	true	false
false	true	false	false	false	false	true	true	true	true	true	true	true
false	true	false	false	false	true	false	true	true	true	true	true	true
false	true	false	false	false	true	true	true	true	true	true	true	true
false	true	false	false	true	false	false	true	false	false	true	true	false
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮
true	true	true	true	true	true	true	false	true	true	false	true	false

5.5 Propositional Theorem Proving

- Determine entailment by **theorem proving**: applying rules of inference directly to the sentences in our knowledge base.
- Some additional concepts related to entailment:
 - Logical equivalence
 - Validity
 - Satisfiability

5.5 Logical equivalence

Two sentences are logically equivalent iff true in same models:

$\alpha \equiv \beta$ if and only if $\alpha \models \beta$ and $\beta \models \alpha$

$(\alpha \wedge \beta) \equiv (\beta \wedge \alpha)$	commutativity of \wedge
$(\alpha \vee \beta) \equiv (\beta \vee \alpha)$	commutativity of \vee
$((\alpha \wedge \beta) \wedge \gamma) \equiv (\alpha \wedge (\beta \wedge \gamma))$	associativity of \wedge
$((\alpha \vee \beta) \vee \gamma) \equiv (\alpha \vee (\beta \vee \gamma))$	associativity of \vee
$\neg(\neg\alpha) \equiv \alpha$	double-negation elimination
$(\alpha \Rightarrow \beta) \equiv (\neg\beta \Rightarrow \neg\alpha)$	contraposition
$(\alpha \Rightarrow \beta) \equiv (\neg\alpha \vee \beta)$	implication elimination
$(\alpha \Leftrightarrow \beta) \equiv ((\alpha \Rightarrow \beta) \wedge (\beta \Rightarrow \alpha))$	biconditional elimination
$\neg(\alpha \wedge \beta) \equiv (\neg\alpha \vee \neg\beta)$	De Morgan
$\neg(\alpha \vee \beta) \equiv (\neg\alpha \wedge \neg\beta)$	De Morgan
$(\alpha \wedge (\beta \vee \gamma)) \equiv ((\alpha \wedge \beta) \vee (\alpha \wedge \gamma))$	distributivity of \wedge over \vee
$(\alpha \vee (\beta \wedge \gamma)) \equiv ((\alpha \vee \beta) \wedge (\alpha \vee \gamma))$	distributivity of \vee over \wedge

Figure 7.11 Standard logical equivalences. The symbols α , β , and γ stand for arbitrary sentences of propositional logic.

5.5 Validity and satisfiability

- A sentence is **valid** if it is true in all models, e.g., True , $A \vee \neg A$, $A \Rightarrow A$
 - Valid sentences are also known as **tautologies**
 - **Validity** is connected to **inference**:
 $KB \models \alpha$ if and only if $(KB \Rightarrow \alpha)$ is valid
- A sentence is **satisfiable** if it is true in some models, e.g., $A \vee B$, C
 - A sentence is **unsatisfiable** if it is true in no models
E.g., $A \wedge \neg A$
 - **Satisfiability** is connected to **inference**:
 $KB \models \alpha$ if and only if $(KB \wedge \neg \alpha)$ is unsatisfiable

5.5 Inference and proofs

Modus Ponens

$$\frac{\alpha \Rightarrow \beta, \quad \alpha}{\beta}$$

And-Elimination

$$\frac{\alpha \wedge \beta}{\alpha}$$

E.g. Wumpus world

R1 : $\neg P_{1,1}$

R2 : $B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})$

R3 : $B_{2,1} \Leftrightarrow (P_{1,1} \vee P_{2,2} \vee P_{3,1})$

R4 : $\neg B_{1,1}$

R5 : $B_{2,1}$

prove $\neg P_{1,2}$

Apply biconditional elimination to R2 to obtain

R6: $(B_{1,1} \Rightarrow (P_{1,2} \vee P_{2,1})) \wedge ((P_{1,2} \vee P_{2,1}) \Rightarrow B_{1,1})$

Apply And-Elimination to R6 to obtain

R7 : $((P_{1,2} \vee P_{2,1}) \Rightarrow B_{1,1})$.

Logical equivalence for contrapositives gives

R8 : $(\neg B_{1,1} \Rightarrow \neg(P_{1,2} \vee P_{2,1}))$.

Apply Modus Ponens with R8 and the percept R4 to obtain

R9: $\neg(P_{1,2} \vee P_{2,1})$.

Apply De Morgan's rule, giving the conclusion

R10 : $\neg P_{1,2} \wedge \neg P_{2,1}$.

5.5 Proof by resolution

- **Conjunctive Normal Form** (CNF—universal)
conjunction of clauses (i.e., disjunctions of literals)
E.g., $(A \vee \neg B) \wedge (B \vee \neg C \vee \neg D)$

Convert $R2 : B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})$ into CNF

1. Eliminate \Leftrightarrow

$$(B_{1,1} \Rightarrow (P_{1,2} \vee P_{2,1})) \wedge ((P_{1,2} \vee P_{2,1}) \Rightarrow B_{1,1}) .$$

2. Eliminate \Rightarrow

$$(\neg B_{1,1} \vee P_{1,2} \vee P_{2,1}) \wedge (\neg(P_{1,2} \vee P_{2,1}) \vee B_{1,1}) .$$

3. Apply logical equivalences

$$(\neg B_{1,1} \vee P_{1,2} \vee P_{2,1}) \wedge ((\neg P_{1,2} \wedge \neg P_{2,1}) \vee B_{1,1}) .$$

4. Apply the distributivity law

$$(\neg B_{1,1} \vee P_{1,2} \vee P_{2,1}) \wedge (\neg P_{1,2} \vee B_{1,1}) \wedge (\neg P_{2,1} \vee B_{1,1}) .$$

5.5 Proof by resolution

- Resolution inference rule (for CNF)

$$\frac{\ell_1 \vee \cdots \vee \ell_k, \quad m_1 \vee \cdots \vee m_n}{\ell_1 \vee \cdots \vee \ell_{i-1} \vee \ell_{i+1} \vee \cdots \vee \ell_k \vee m_1 \vee \cdots \vee m_{j-1} \vee m_{j+1} \vee \cdots \vee m_n}$$

where ℓ_i and m_j are complementary literals (i.e., one is the negation of the other).

5.5 Proof by resolution

- Inference procedures based on a resolution algorithm uses the principle of proof by contradiction

$KB \models \alpha$ if and only if $(KB \wedge \neg\alpha)$ is unsatisfiable

- Steps:
 1. $(KB \wedge \neg\alpha)$ is converted into CNF
 2. The resolution rule is applied to the resulting clauses, a new clause is added to the set if it is not already present
 3. The process continues until one of two things happens:
 - i. no new clauses that can be added, in which case KB does not entail α ;
 - ii. two clauses resolve to yield the empty clause (\sim False), in which case KB entails α .

5.5 Proof by resolution

E.g. Wumpus world

$R_1 : \neg P_{1,1}.$

$R_2 : B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1}).$

$R_3 : B_{2,1} \Leftrightarrow (P_{1,1} \vee P_{2,2} \vee P_{3,1}).$

$R_4 : \neg B_{1,1}.$

$R_5 : B_{2,1}.$

Prove $\neg P_{1,2}$ by resolution

5.5 Proof by resolution

$KB = R2 \wedge R4 = (B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})) \wedge \neg B_{1,1}$

α as $\neg P_{1,2}$

1. convert $(KB \wedge \neg \alpha)$ to CNF

$(\neg P_{1,2} \vee B_{1,1}) \wedge (\neg B_{1,1} \vee P_{1,2} \vee P_{2,1})$
 $\wedge (\neg P_{2,1} \vee B_{1,1}) \wedge \neg B_{1,1} \wedge \neg P_{1,2}$

2. resolve pairs

$(\neg P_{1,2} \vee B_{1,1}), (\neg B_{1,1} \vee P_{1,2} \vee P_{2,1}): P_{2,1}$

$(\neg P_{1,2} \vee B_{1,1}), \neg B_{1,1}: \neg P_{1,2}$

$(\neg P_{2,1} \vee B_{1,1}), (\neg B_{1,1} \vee P_{1,2} \vee P_{2,1}): P_{1,2}$

$(\neg P_{2,1} \vee B_{1,1}), \neg B_{1,1}: P_{2,1}$

3. resolve pairs

$\neg P_{1,2}, P_{1,2}$: empty

Result: $KB \models \neg P_{1,2}$

5.5 Proof by resolution

Proof by contradiction, i.e., show $KB \wedge \neg\alpha$ unsatisfiable

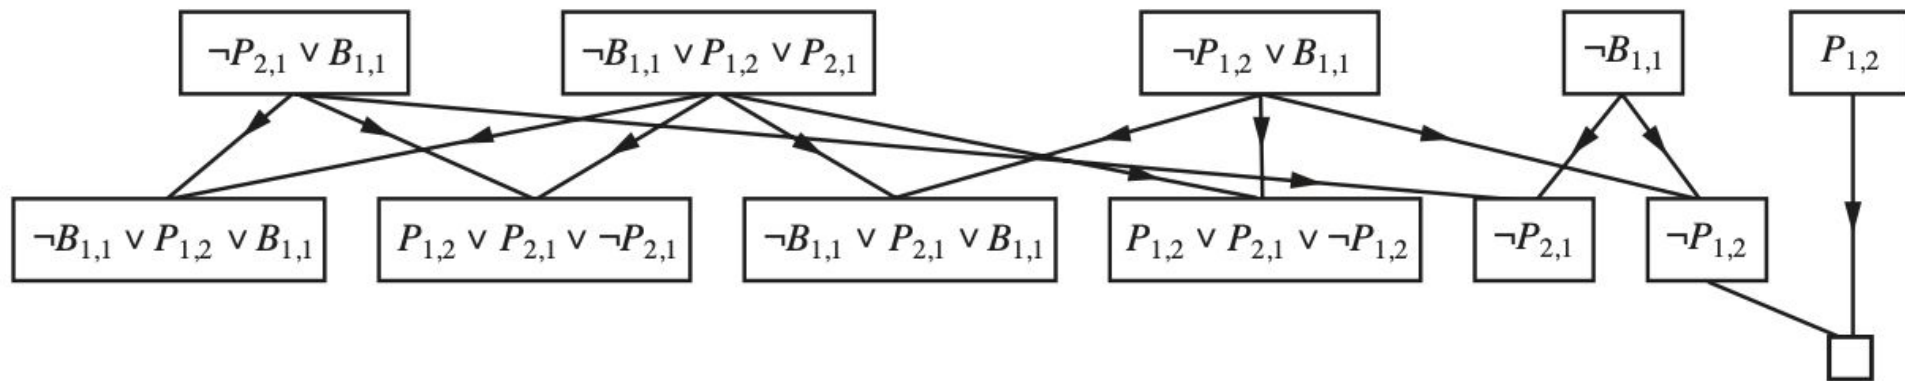
```
function PL-RESOLUTION( $KB, \alpha$ ) returns true or false
  inputs:  $KB$ , the knowledge base, a sentence in propositional logic
            $\alpha$ , the query, a sentence in propositional logic

   $clauses \leftarrow$  the set of clauses in the CNF representation of  $KB \wedge \neg\alpha$ 
   $new \leftarrow \{ \}$ 
  loop do
    for each  $C_i, C_j$  in  $clauses$  do
       $resolvents \leftarrow$  PL-RESOLVE( $C_i, C_j$ )
      if  $resolvents$  contains the empty clause then return true
       $new \leftarrow new \cup resolvents$ 
  if  $new \subseteq clauses$  then return false
   $clauses \leftarrow clauses \cup new$ 
```


5.5 Proof by resolution

$$KB = (B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})) \wedge \neg B_{1,1}$$

$$\alpha = \neg P_{1,2}$$



5.5 Horn clauses and definite clauses

- **Definite clause:** a disjunction of literals of which **exactly one** is positive.
E.g., $(\neg L_{1,1} \vee \neg \text{Breeze} \vee B_{1,1})$ is a definite clause
- **Horn clause:** a disjunction of literals of which **at most one** is positive
- **Goal clauses:** clauses with **no positive** literals

$$CNFSentence \rightarrow Clause_1 \wedge \cdots \wedge Clause_n$$

$$Clause \rightarrow Literal_1 \vee \cdots \vee Literal_m$$

$$Literal \rightarrow Symbol \mid \neg Symbol$$

$$Symbol \rightarrow P \mid Q \mid R \mid \dots$$

$$HornClauseForm \rightarrow DefiniteClauseForm \mid GoalClauseForm$$

$$DefiniteClauseForm \rightarrow (Symbol_1 \wedge \cdots \wedge Symbol_l) \Rightarrow Symbol$$

$$GoalClauseForm \rightarrow (Symbol_1 \wedge \cdots \wedge Symbol_l) \Rightarrow False$$

Figure 7.14 A grammar for conjunctive normal form, Horn clauses, and definite clauses.

5.5 Forward chaining

Idea: fire any rule whose premises are satisfied in the KB,
add its conclusion to the KB, until query is found or no further inferences can be made.

```
function PL-FC-ENTAILS?(KB, q) returns true or false
  inputs: KB, the knowledge base, a set of propositional Horn clauses
           q, the query, a proposition symbol
  local variables: count, a table, indexed by clause, initially the number of premises
                    inferred, a table, indexed by symbol, each entry initially false
                    agenda, a list of symbols, initially the symbols known in KB

  while agenda is not empty do
    p ← POP(agenda)
    unless inferred[p] do
      inferred[p] ← true
      for each Horn clause c in whose premise p appears do
        decrement count[c]
        if count[c] = 0 then do
          if HEAD[c] = q then return true
          PUSH(HEAD[c], agenda)

  return false
```

5.5 Forward chaining

- In AND–OR graphs,
 - multiple links joined by an arc indicate a conjunction
 - multiple links without an arc indicate a disjunction
- How the graphs work:
 - The known leaves are set, inference propagates up the graph as far as possible.
 - Where a conjunction appears, the propagation waits until all the conjuncts are known before proceeding.

$$P \Rightarrow Q$$

$$L \wedge M \Rightarrow P$$


$$B \wedge L \Rightarrow M$$

$$A \wedge P \Rightarrow L$$

$$A \wedge B \Rightarrow L$$

A

B



5.5 Forward chaining

$$P \Rightarrow Q$$

$$L \wedge M \Rightarrow P$$



$$B \wedge L \Rightarrow M$$

$$A \wedge P \Rightarrow L$$

$$A \wedge B \Rightarrow L$$

A

B



5.5 Forward chaining

$$P \Rightarrow Q$$

$$L \wedge M \Rightarrow P$$



$$B \wedge L \Rightarrow M$$

$$A \wedge P \Rightarrow L$$

$$A \wedge B \Rightarrow L$$

A

B



5.5 Forward chaining

$$P \Rightarrow Q$$

$$L \wedge M \Rightarrow P$$


$$B \wedge L \Rightarrow M$$

$$A \wedge P \Rightarrow L$$

$$A \wedge B \Rightarrow L$$

A

B



5.5 Forward chaining

$P \Rightarrow Q$

$L \wedge M \Rightarrow P$



$B \wedge L \Rightarrow M$

$A \wedge P \Rightarrow L$

$A \wedge B \Rightarrow L$

A




B






5.5 Backward chaining

- Idea: work backwards from the query q to prove q by BC,
 - check if q is known already, or
 - prove by BC all premises of some rule concluding q
- Avoid loops: check if new subgoal is already on the goal stack
- Avoid repeated work: check if new subgoal
 - 1) has already been proved true, or
 - 2) has already failed




5.5 Backward chaining





5.5 Backward chaining



5.5 Backward chaining



5.5 Backward chaining



5.5 Forward vs. backward chaining

- FC is data-driven, cf. automatic, unconscious processing,
 - e.g., object recognition, routine decisions
 - May do lots of work that is irrelevant to the goal
- BC is goal-driven, appropriate for problem-solving,
 - e.g., Where are my keys? How do I get into a PhD program?